

OMNIRA ARCHITECTURE SERIES

EXECUTIVE INTELLIGENCE

Canonical Architecture and Operating Doctrine

PROFESSIONAL EDITION · VERSION 1.0

Omnira — Executive Intelligence

Canonical Architecture and Operating Doctrine

TITLE	Omnira — Executive Intelligence
SUBTITLE	Canonical Architecture and Operating Doctrine
EDITION	Professional Edition v1.0 (derivative of Canonical Edition v1.0)
VERSION	1.0
CANONICAL STATUS	Approved and locked
APPROVAL DATE	2026-07-08
DOCUMENT OWNER	André Hultgren
SOURCE LINEAGE	Phase 1.2 Editorial Review Draft
CANONICAL SHA-256	ee85a1a0...0fa555b8

Professional Edition — visual derivative. Canonical text is reproduced verbatim and remains authoritative.

Document Control

Document	Omnira — Executive Intelligence — Professional Edition v1.0
Canonical source	Executive Intelligence — Canonical Edition v1.0
Version	1.0
Canonical status	Approved and locked
Approval date	2026-07-08
Document owner	André Hultgren

Canonical Doctrine Notice

Version 1.0 is the approved canonical target architecture for Omnira Executive Intelligence. This book describes both the Stage 1 scope and the future target architecture.

Canonical target architecture does not mean that every described capability is implemented. The actual repository, schema, runtime, and deployment state is the authoritative source for what exists in production. Atlas and future implementations must distinguish between canonical target architecture, implemented, partially implemented, planned, not implemented, and deprecated.

This book provides architectural knowledge but grants no automatic execution authority. No document, memory, knowledge object, graph object, or model output may create authority without valid governance.

Table of Contents

Live navigation · Part grouping · clickable internal links

PART V — GOVERNANCE, DAMAGE & AUTONOMY

Chapter 16 — Governance & Policy Engine	6
Chapter 17 — Damage Boundary	10
Diagram — Omnira Intelligence Layer Stack	11
Diagram — Executive · Manager · Workforce	12
Diagram — Damage Severity & Boundary States	13
Diagram — Stage 1 vs Future Target	14

Non-canonical navigational grouping added only in the Professional Edition.

PART V

GOVERNANCE, DAMAGE & AUTONOMY

Chapters 16 – 19

Non-canonical navigational grouping added only in the Professional Edition.

16

Governance & Policy Engine

16.1 Purpose of This Chapter

Executive Intelligence must be able to reason broadly.

It must not be able to act broadly without authority.

A recommendation may be strategically intelligent.

A plan may create revenue.

A workflow may be technically capable.

An agent may have access to the required tool.

None of these facts alone mean that the action is allowed.

16.1 Purpose of This Chapter

Executive Intelligence must be able to reason broadly.

It must not be able to act broadly without authority.

A recommendation may be strategically intelligent.

A plan may create revenue.

A workflow may be technically capable.

An agent may have access to the required tool.

None of these facts alone mean that the action is allowed.

Omnira therefore requires a Governance & Policy Engine that determines:

- Is this action permitted?
- Who has authority?
- Which policy applies?
- Is approval required?
- Does an autonomy license cover the action?
- Does the action cross the Damage Boundary?
- Does project scope match?
- Does a higher-level rule prohibit it?

25.1 Purpose of This Chapter

Executive Intelligence must understand whether the organization is producing the intended outcomes.

It must be able to evaluate:

- Revenue.
- Customer behavior.
- Agent quality.
- Founder burden.
- Cost.
- Content performance.
- Mission progress.
- Strategic execution.
- Growth.
- Workflow reliability.
- Autonomy outcomes.
- Return on investment.

However, raw metrics do not explain themselves.

A metric may improve while the business becomes weaker.

A workflow may produce more output while creating less value.

Revenue may increase because of a temporary effect.

Cost may fall because important work stopped.

An agent may appear productive because difficult cases are escalated elsewhere.

A dashboard may show activity without showing progress.

The purpose is evidence-based leadership without metric-driven blindness.

Callout System

CANONICAL DOCTRINE

This book provides architectural knowledge but grants no automatic execution authority. No document, memory, knowledge object, graph object, or model output may create authority without valid governance.

Front matter · Canonical Doctrine Notice

STAGE 1 BOUNDARY

Implementation should occur incrementally. Stage 1 should be recommendation-first and human-authorized, establishing a usable Executive foundation without pretending full autonomy exists.

Front matter · Implementation Scope and Maturity

FUTURE TARGET ARCHITECTURE

Future systems remain architectural contracts until separately designed, approved, implemented, validated, and governed.

Front matter · Implementation Scope and Maturity

EXECUTIVE DECISION

Executive Intelligence must be able to reason broadly. It must not be able to act broadly without authority.

Ch 16 §16.1

DAMAGE BOUNDARY

The purpose is to prevent meaningful harm from being created without appropriate human authority and control.

Ch 17 §17.1

AUTONOMY LICENSE

Autonomy may exist only below a defined Damage Boundary or within a licensed controlled zone.

Ch 17

ANTI-PRINCIPLE

If Executive Intelligence performs specialist work directly, it becomes a bottleneck.

Ch 3 §3.3

EXAMPLE

A metric may improve while the business becomes weaker.

Ch 25 §25.1

All callout bodies are verbatim canonical text; labels are design containers only.


17.100 Damage Severity

Class	Level	Canonical response
D0	Negligible	Normal autonomous handling may be appropriate.
D1	Limited	Logging and local correction may be sufficient.
D2	Material	Human review or explicit bounded authority is generally required.
D3	Severe	Immediate containment and senior approval are required.
D4	Critical or Systemic	Crisis Mode should activate.

This classification is distinct from policy violation severity, though related.

Canonical responses drawn verbatim from Ch 17 §§17.101-17.105.

Omnira Intelligence Layer Stack



Canonical position model – Chapter 2 §§2.2-2.3.

Solid = authority & direction · dashed = communication / resource selection · gold loop = evidence feedback

Executive · Manager · Workforce



- Executive leadership (gold)
- Manager coordination (navy)
- Workforce execution (grey)
- Collaboration, not authority

Damage Severity & Boundary States


PANEL A — DAMAGE SEVERITY	PANEL B — BOUNDARY STATES
D0 — Negligible Normal autonomous handling may be appropriate.	Below Boundary The action may proceed within valid authority.
D1 — Limited Logging and local correction may be sufficient.	Near Boundary The action may proceed only with additional controls.
D2 — Material Human review or explicit bounded authority is generally required.	Crosses Boundary The action requires explicit human authority or a valid high-confidence mandate.
D3 — Severe Immediate containment and senior approval are required.	Prohibited Regardless of Approval Some actions may remain prohibited because they violate: Law, Customer ownership, Constitutional rules, Security requirements, Non-negotiable safety boundaries.
D4 — Critical or Systemic Crisis Mode should activate.	

Two distinct canonical classification systems (Ch 17 §§17.100-17.119). No one-to-one mapping between severity classes and boundary states is implied.

"Prohibited Regardless of Approval" is a separate boundary state, not a severity class. D4 renders: Crisis Mode should activate.

- Autonomous / low
- Approval / added controls
- Crisis / prohibited

Stage 1 vs Future Target Architecture



Source: Front matter – Implementation Scope and Maturity (verbatim canonical scope lists).

- Stage 1 scope (solid)
- Future target (dashed)

Chapter transition & navigation

This closing proof page demonstrates the running header, footer, page number, and chapter track that carry the reader across the full book. In production each chapter ends with a thin Petrol rule and the next chapter opens on a dark feature page.

Next · Chapter 18 — Autonomy Licensing Model

PART V — GOVERNANCE, DAMAGE & AUTONOMY · Chapters 16–19